

# Project1

November 11, 2024

## 1 Project 1

The following packages are used in this project

```
[1]: import numpy as np
from scipy.sparse import csr_matrix
import scipy.sparse.linalg as la
import matplotlib.pyplot as plt
import seaborn as sns
import scipy
import time
```

1. Write a function that loads this file and returns the matrix X and vector y. How many malicious data points are there? What can you say about the sparsity of the data? Do you think it makes sense to use one-hot-coding for some of the columns?

we use the following code to loads the data file, returns the matrix X and vector y, and count the number of malicious data points:

```
[2]: class DataLoader:
    def __init__(self, file_path= r".\data\data.csv"):
        self.X = None
        self.y = None
        self.num = 0

        self._read(file_path)

    def _read(self, file_path):
        ori_data = np.loadtxt(file_path, skiprows=1, delimiter=',', dtype=int)
        self.X = ori_data[:, :-1]
        self.y = ori_data[:, -1]
        self.y[self.y == 0] = -1
        self.num = self.X.shape[0]

    def split(self, ratio = 0.7):
        # the number of training data
        train_num = int(self.num * ratio)
```

```

        # shuffling the original data
        rand_idx = np.arange(self.num)
        np.random.shuffle(rand_idx)
        X = self.X[rand_idx]
        y = self.y[rand_idx]

        # split the data set
        X_train, X_test = X[:train_num], X[train_num:]
        y_train, y_test = y[:train_num], y[train_num:]

        return X_train, X_test, y_train, y_test

mydata = DataLoader(r".\data\data.csv")
print(f'there are {(mydata.y == 1).sum()} \
malicious data points in the total {mydata.num} data points')

```

there are 14700 malicious data points in the total 29332 data points

From summing the ones in the data, we get that there are 14700 malicious data points in the total 29332 data points.

The dataset is highly sparse, as most feature values are zero.

For onw-hot encoding, it is reasonable to apply it when the features represent categorical data without intrinsic ordering. However, for binary features in our project, one-hot encoding may be redundant unless further categorization is needed.

2. Write a function that splits the data into a training and test set according to some fraction  $0 < r < 1$ . Make sure to use randomization; that is, it should not be the case that the training set consists of the first data points and the test set of the remaining data points. Your function should return matrices  $X$  train and  $X$  test and vectors  $y$  train and  $y$  test.

We used the `split` method from the `DataLoader` class shown in **Question 1** to divide the dataset according to the given ratio, where shuffling is implemented before splitting to avoid the case that the training set consists of the first data points and the test set of the remaining data points.

```

[3]: np.random.seed(1234)
      X_train, X_test, y_train, y_test = mydata.split(ratio = 0.5)

```

3. Write a function that, given the matrix  $X$ , the vector  $y$ , and a weight vector  $w$  defining a hyperplane, returns the number of correctly classified points. Verify that the output makes sense for random weight vectors.

We used a randomly generated weight vector `w_init` to define a hyperplane. The following function is used to compute predictions and the number of correctly classified points for the given input matrix  $X$ , output  $y$  and weight vector  $w$ .

```

[4]: def evaluator(X: np.ndarray, y: np.ndarray, w: np.ndarray):
      y_pred = X.dot(w)
      y_pred[y_pred < 0] = -1

```

```

y_pred[y_pred >= 0] = 1
return (y_pred == y).sum()

np.random.seed(1234)
w_init = np.random.uniform(low = -1, high = 1, size=(mydata.X.shape[1]))
pred_corr_num = evaluator(mydata.X, mydata.y, w_init)
print(f'{pred_corr_num} data points in total {mydata.num} samples are correct_
↳using random generated weights')

```

11754 data points in total 29332 samples are correct using random generated weights

4. Consider the cost function for logistic regression as defined in the lectures. Write down a symbolic formula for the gradient of this function.

The cost function  $J(w)$  for logistic regression with regularization is defined as:

$$J(w) = \sum_{i=1}^n L(y_i x_i^T w) + \frac{\lambda}{2} \|w\|^2$$

where:

$$L(s) = \log(1 + e^{-s})$$

The gradient of the loss term for a single data point  $x_i$  with respect to  $w$  is:

$$\nabla_w L(y_i x_i^T w) = -\frac{y_i x_i}{1 + e^{y_i x_i^T w}}$$

Therefore, the gradient of the total loss term is:

$$\sum_{i=1}^n \nabla_w L(y_i x_i^T w) = -\sum_{i=1}^n \frac{y_i x_i}{1 + e^{y_i x_i^T w}}$$

The gradient of the regularization term  $\frac{\lambda}{2} \|w\|^2$  with respect to  $w$  is:

$$\nabla_w \left( \frac{\lambda}{2} \|w\|^2 \right) = \lambda w$$

Thus, the gradient of  $J(w)$  with respect to  $w$  is:

$$\nabla_w J(w) = -\sum_{i=1}^n \frac{y_i x_i}{1 + e^{y_i x_i^T w}} + \lambda w$$

5. • Write a straightforward implementation for logistic regression using gradient descent with a fixed step size  $\alpha$ . Your function should take as arguments the data matrix  $X$  and data vector  $y$ , the step size  $\alpha$ , the regularization constant

$\lambda$ , and an integer  $K$  indicating the number of gradient descent steps. The function should return a weight vector  $w$ .

- Experiment with the hyperparameters, using dense and sparse linear algebra, on random splits of training and test data sets. (If you know about writing allocation-free code, you can also experiment with this.)
- Given for instance a 50/50 split between test and training data, what is the best classification performance you can obtain on the test set?

In this logistic regression implementation, the constants  $\alpha = 0.01$  (step size),  $\lambda = 0.01$  (regularization constant), and  $K = 100$  (gradient descent steps) are chosen initially to test our code.

We implement our code in 3 different ways: 1. The `logistic_regression_entrywise` method calculates the gradient individually for each data point in every iteration.

2. The `logistic_regression_vectorize` method utilizes vectorizing and broadcasting to avoid explicitly calculating the gradient for each data point individually, show as below:

- $s = y * (X @ w)$
- $z = y / (1 + \text{np.exp}(s))$
- $w_{\text{grad}} = - (X.T @ z) + \text{lambda\_} * w$

3. The `logistic_regression_sparse` method is further optimized by converting training data to a CSR type sparse matrix and using sparse linear algebra, due to the sparsity in our training data.

The code implement above is shown as below:

```
[5]: class LogisticClassifier:
    def __init__(self, X_train, y_train, w_init):
        self.X_train = X_train
        self.y_train = y_train
        self.w_init = w_init
        self.w = None

    def logistic_regression(self, K, alpha, lambda_reg, flag = 0):
        if flag == 0:
            self.logistic_regression_entrywise(K, alpha, lambda_reg)
        if flag == 1:
            self.logistic_regression_vectorize(K, alpha, lambda_reg)
        if flag == 2:
            self.logistic_regression_sparse(K, alpha, lambda_reg)

    def logistic_regression_entrywise(self, K, alpha, lambda_reg):
        self.w = self.w_init
        for k in range(K):
            w_grad = lambda_reg * self.w
            for i in range(self.X_train.shape[0]):
                # if k == 5:
                #     print(i)
                w_grad -= (self.y_train[i] * self.X_train[i]) \
                    / (1 + np.exp(self.y_train[i] * self.X_train[i].dot(self.w)))
```

```

        self.w -= w_grad * alpha
    return self.w

def logistic_regression_vectorize(self, K, alpha, lambda_reg):
    self.w = self.w_init
    for _ in range(K):
        s = self.y_train[:,None] * self.X_train @ self.w
        z = self.y_train / (1 + np.exp(s))
        w_grad = - self.X_train.T @ z + lambda_reg * self.w
        self.w -= w_grad * alpha
    return self.w

def logistic_regression_sparse(self, K, alpha, lambda_reg):
    self.w = self.w_init
    self.X_train = csr_matrix(self.X_train)
    for _ in range(K):
        s_product = self.X_train @ self.w
        s = self.y_train * s_product
        z = self.y_train / (1 + np.exp(s))
        w_grad = - self.X_train.T @ z + lambda_reg * self.w
        self.w -= w_grad * alpha
    return self.w

def predict(self, X_test, y_test):
    w = self.w
    y_true = evaluator(X_test, y_test, w)
    acc = y_true / y_test.shape[0]
    return acc

```

```

[6]: log_reg = LogisticClassifier(X_train, y_train, w_init)

start = time.time()
log_reg.logistic_regression(K=100, alpha=0.01, lambda_reg=0.01, flag= 0)
accuracy = log_reg.predict(X_test, y_test)
end = time.time()
print(f'Entry-Wise Implement:\n Accuracy on the test set: {accuracy * 100:.
    ↪2f}%, the time cost is {end-start}s\n')

start = time.time()
log_reg.logistic_regression(K=100, alpha=0.01, lambda_reg=0.01, flag=1)
accuracy = log_reg.predict(X_test, y_test)
end = time.time()
print(f'Dense Implement:\nAccuracy on the test set: {accuracy * 100:.2f}%, the
    ↪time cost is {end-start}s\n')

start = time.time()
log_reg.logistic_regression(K=100, alpha=0.01, lambda_reg=0.01, flag=2)

```

```

accuracy = log_reg.predict(X_test, y_test)
end = time.time()
print(f'Sparse Implement:\nAccuracy on the test set: {accuracy * 100:.2f}%, the_
↳time cost is {end - start}s\n')

```

Entry-Wise Implement:

Accuracy on the test set: 94.81%, the time cost is 14.146091222763062s

Dense Implement:

Accuracy on the test set: 94.74%, the time cost is 2.2567808628082275s

Sparse Implement:

Accuracy on the test set: 95.30%, the time cost is 0.31095409393310547s

From the above experiment, we have 1. The three implements achieve similar accuracy with different running times.

2. The `logistic_regression_entrywise` method is quite slow for large datasets because of the iterative sample-by-sample gradient computation, taking approximately 14.15 seconds for 100 iteration.
3. The `logistic_regression_vectorize` method reduces computation time to roughly 2.26 second for 100 iteration.
4. The `logistic_regression_sparse` method is the fastest with a rapid runtime of 0.31 seconds for 100 iteration, making it ideal for large, sparse datasets by reducing unnecessary calculations.

By experimenting on hyperparameters, the best classification performance we can obtain the best accuracy of 95.58% in 1000 iterations on the test set with a 50/50 split with  $\alpha = 0.03$  and  $\lambda = 0.1$  among hyperparameters we tested.

```

[7]: best_alpha = None
best_lambda_reg = None
best_acc = 0
for alpha in [0.1,0.03,0.01,0.003,0.001]:
    for lambda_reg in [0.1,0.01,0.001,0.0001]:
        log_reg.logistic_regression(K=1000, alpha=alpha, lambda_reg=lambda_reg,
↳flag=2)
        accuracy = log_reg.predict(X_test, y_test)
        if accuracy >= best_acc:
            best_acc = accuracy
            best_alpha = alpha
            best_lambda_reg = lambda_reg

print(f'best_acc = {best_acc}, best_alpha = {best_alpha}, best_lambda_reg =_
↳{best_lambda_reg}')

```

C:\Users\11787\AppData\Local\Temp\ipykernel\_6352\3012149010.py:43:

RuntimeWarning: overflow encountered in exp

```

z = self.y_train / (1 + np.exp(s))

best_acc = 0.9558161734624301, best_alpha = 0.03, best_lambda_reg = 0.1

```

- Download the file data2.csv. This is the same as the previous data file, except that 2000 fake data points have been appended to the data set. Use the singular value transform as explained in the lecture to detect and remove most of these outliers without removing too many other data points. (For the singular value transform you can use a library function, you do not have to implement this yourself.)

We use the following function to whitening our data:

```

[8]: def data_whitening(X, k):
    # Centering
    M = X - X.mean(axis=0)

    # SVD Decomposition
    U, S, _ = la.svds(M, k=k)

    # Whitening
    Xw = U[:,::-1] @ scipy.sparse.diags(S[:,::-1])
    Xw = Xw / np.std(Xw, axis=0)[None,:]

    return Xw

```

Then we compute the norm of whitening data to detect the outliers in our data, Within the noise\_removal function, we employed the `scipy.sparse.linalg.svds` method to compute the top k singular values and their corresponding singular vectors:

```

[9]: def noise_removal(X, k, percentile = 0.9, vis_flag = False):
    # Whitening
    Xw = data_whitening(X, k)

    norm_list = np.array([np.linalg.norm(Xw[i, :]) for i in range(Xw.shape[0])])
    outlier_mask = norm_list >= np.quantile(norm_list, percentile)

    return outlier_mask

```

Let - TP (True positive) = number of fake samples be detected - FN (False Negative) = number of fake samples be identified as true data points - TN (True Negative) = number of true data points be correctly identified - FP (False Positive) = number of true data point be wrongly identified as fake samples.

We use Precision, Recall and F1-score to assess the performance of our method

```

[10]: def assess_performance(outlier_mask, vis_flag = False):
    confusion_matrix, (P, R), F1 = F1_score(outlier_mask)
    acc = (confusion_matrix[0] + confusion_matrix[2]) /\
          (confusion_matrix[0] + confusion_matrix[1] + confusion_matrix[2]
    ↪ + confusion_matrix[3])

```

```

if vis_flag:
    visualize_confusion(*confusion_matrix)
    print(f'Accuracy: {acc:.4f}')
    print(f"Precision: {P:.4f}")
    print(f"Recall: {R:.4f}")
    print(f"F1-score: {F1:.4f}")

return F1, (P, R), acc

def F1_score(outlier_mask, outlier_num = 2000):
    """
    The outlier is in the last #outlier_num rows.
    """

    total_num = len(outlier_mask)

    TP = (outlier_mask[-outlier_num:].sum())
    FP = (outlier_mask[:-outlier_num].sum())

    TN = total_num - outlier_num - FP
    FN = outlier_num - TP

    Precision = TP / (TP + FP)
    Recall = TP / (TP + FN)
    F1_score = 2 * (Precision * Recall) / (Precision + Recall) if (Precision +
↪Recall) > 0 else 0

    return (TP, FP, TN, FN), (Precision, Recall), F1_score

def visualize_confusion(TP, FP, TN, FN):
    confusion_matrix = np.array([[TN, FP],
                                [FN, TP]])

    categories = ['Normal', 'Outlier']

    # plot confusion matrix
    plt.figure(figsize=(6, 4))
    sns.heatmap(confusion_matrix, annot=True, fmt='d', cmap='Blues',
                xticklabels=categories, yticklabels=categories)
    plt.xlabel('Prediction')
    plt.ylabel('Actual')
    plt.title('Confusion Matrix')
    plt.show()

```

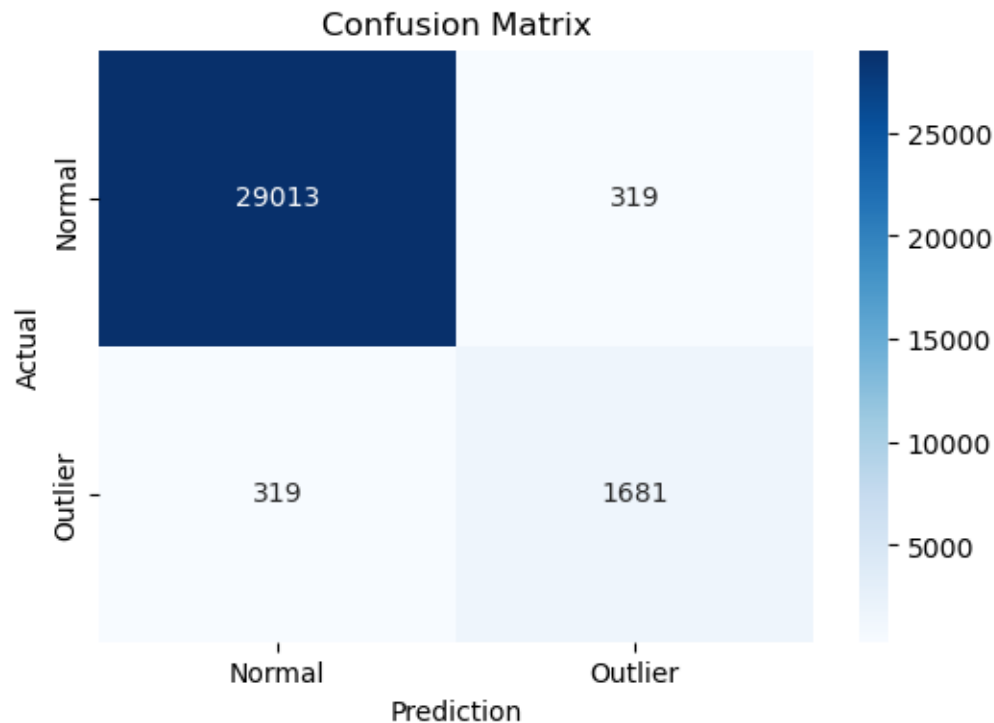
The we have the following result:



```
[11]: # read data
my_data = DataLoader(r".\data\data2.csv")
X = my_data.X.astype('float64')

# Outlier rate
q = 1 - 2000 / X.shape[0]

# data whitening & noise removal
F1, _, _ = assess_performance(noise_removal(X, 81, q), vis_flag=True)
```



Accuracy: 0.9796  
Precision: 0.8405  
Recall: 0.8405  
F1-score: 0.8405

In the plot of confusion matrix, we found that: - The diagonal elements (29,013 normal points and 1,681 outliers) indicate correctly classified instances.

- Off-diagonal elements (319 misclassified as normal and 319 as outliers) indicate errors
- Majority of outliers was removed and only a few true data point was misclassifieds.

By adjusting the threshold for classifying a data point as an outlier or not, we obtain different classification results, focusing either on detecting more true positives (TP) or on detecting true positives with greater accuracy. The precision-recall curve is shown below

```

[14]: def plot_PRC(P_list, R_list):
    # data clean
    valid_indices = (~np.isnan(P_list)) & (~np.isnan(R_list)) & (~np.
    ↪ isinf(P_list)) & (~np.isinf(R_list))
    P_array = P_list[valid_indices]
    R_array = R_list[valid_indices]

    # sort data
    sorted_indices = np.argsort(R_array)
    R_sorted = R_array[sorted_indices]
    P_sorted = P_array[sorted_indices]

    # plot data
    plt.figure(figsize=(8, 6))
    plt.plot(R_sorted, P_sorted, linestyle='--', color='b')
    plt.xlabel('Recall', fontsize=12)
    plt.ylabel('Precision', fontsize=12)
    plt.title('Precision-Recall Curve', fontsize=14)
    plt.grid(True)
    plt.show()

P_list = []
R_list = []
for q in np.arange(0, 1, 0.01):
    print(f"Progress: {q * 100:.0f}% (q = {q:.2f})", end='\r')
    _, (P, R), _ = assess_performance(noise_removal(X, 81, q), vis_flag=False)
    P_list.append(P)
    R_list.append(R)

# plot PR-Curve
plot_PRC(np.array(P_list), np.array(R_list))

```

Progress: 99% (q = 0.99)

